



# P1.2 course:

## DAY 2: Introduction to modern CPUs

Stefano Cozzini

CNR/IOM and eXact-lab srl



Scuola Internazionale Superiore  
di Studi Avanzati



## Outline

- Moore law
- Components of a modern CPU
- Von Neumann architecture
- Modern features of CPU and cores
  - SIMD
  - Pipelines
  - superscalar processing and ILP
- Survey of modern CPUs for HPC
- Final consideration

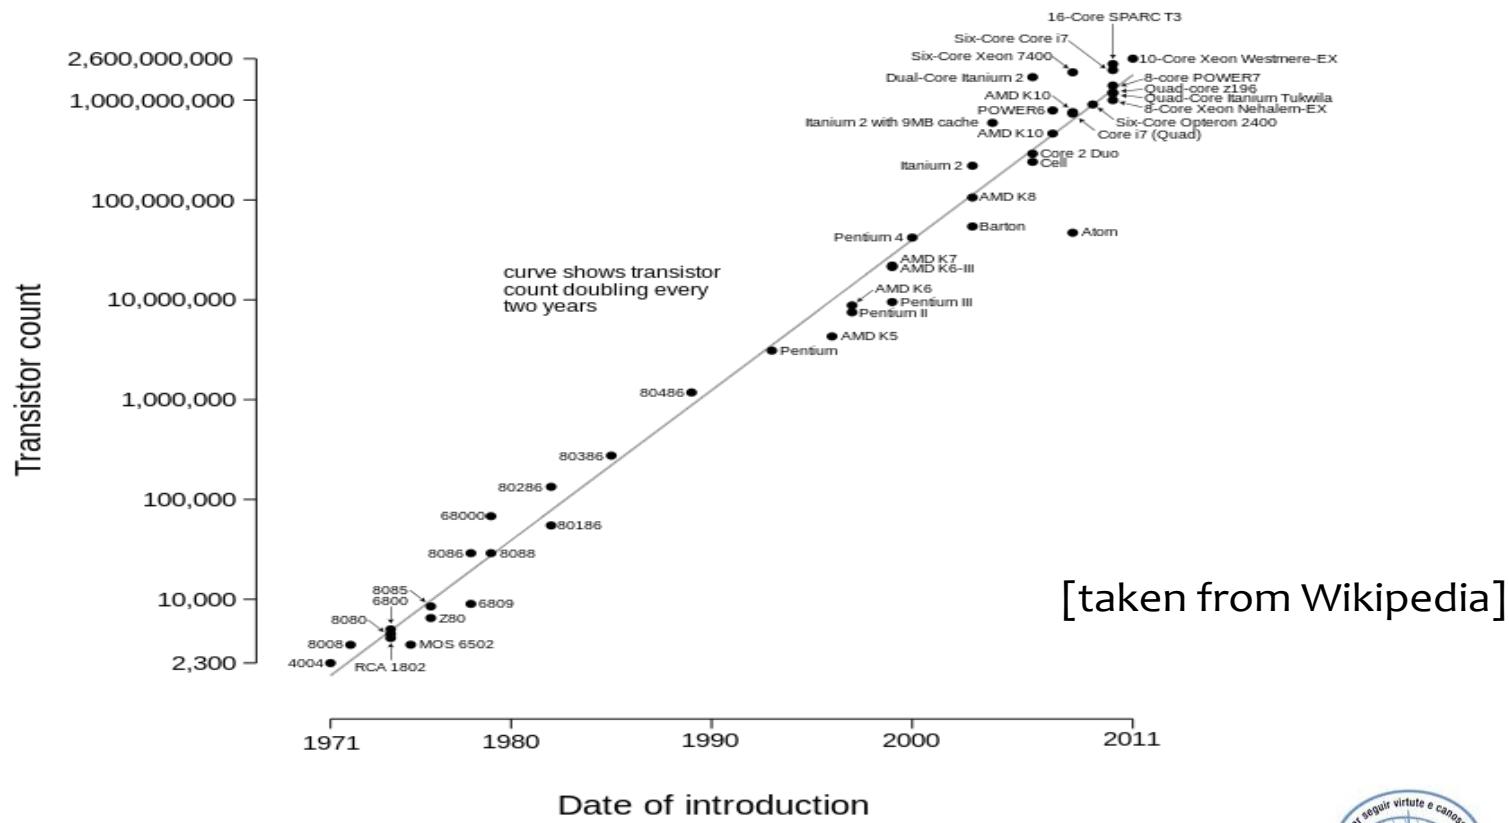
## Moore's law

- The number of transistors in a dense integrated circuit doubles approximately every two years.  
[ Gordon Moore, the co-founder of Intel and Fairchild Semiconductor, 1965 ]
- The period is often quoted as 18 months because of Intel executive David House, who predicted that **chip performance would double every 18 months** (being a combination of the effect of more transistors and the transistors being faster)

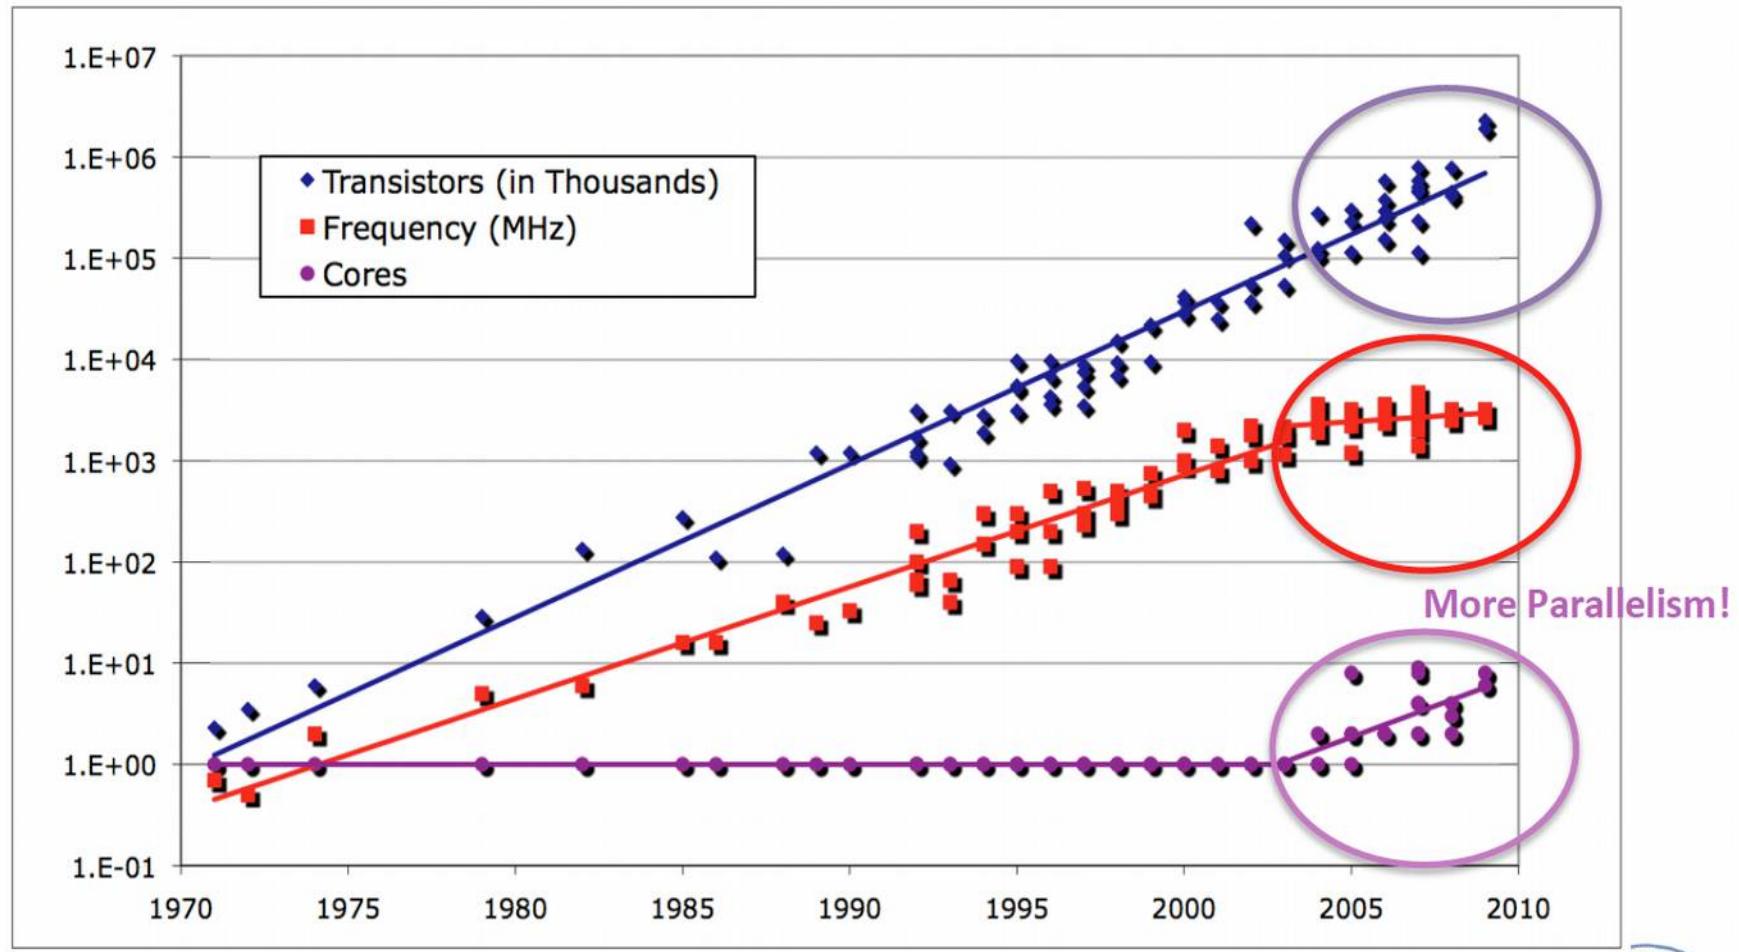
[adapted from Wikipedia]

## Moore's law:

Microprocessor Transistor Counts 1971-2011 & Moore's Law



## CPUs and Moore law



## Observations

- Silicon technology is slowing down Moore's law
- End of the law in 5 years ?
- Exploitation of Moore law means parallelism at all level..
- We explore it at core/cpu level today..

a picture from Intel..



**"Parallelism for Everyone"**

Parallelism changes the game

- A large percentage of people who provide applications are going to have to care about parallelism in order to match the capabilities of their competitors.

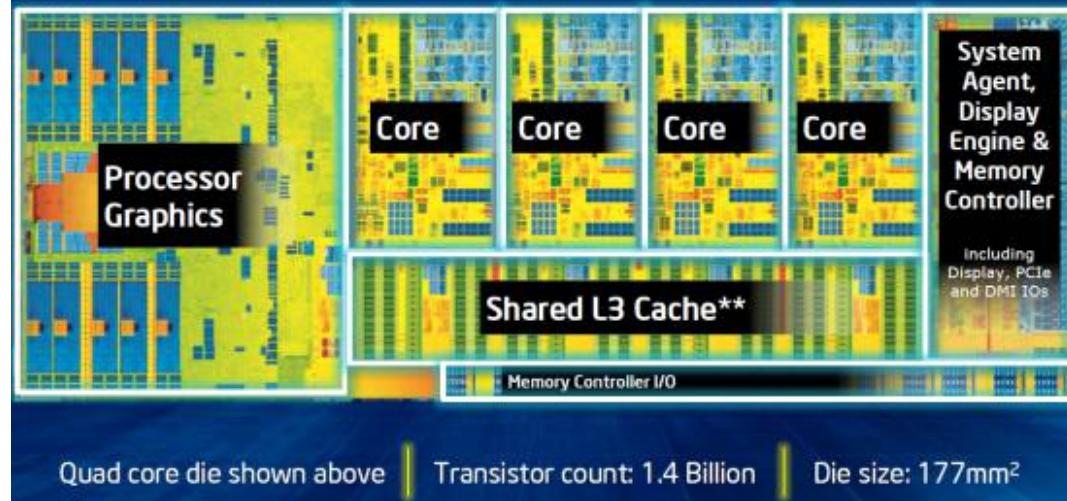
## Something to read on Moore's law

- <https://www.technologyreview.com/s/601441/moores-law-is-dead-now-what/>
- <https://www.kth.se/social/upload/507d1d3af27654019000002/Moore's%20law.pdf>  
(available on the github repository)

## What does a CPU look like?

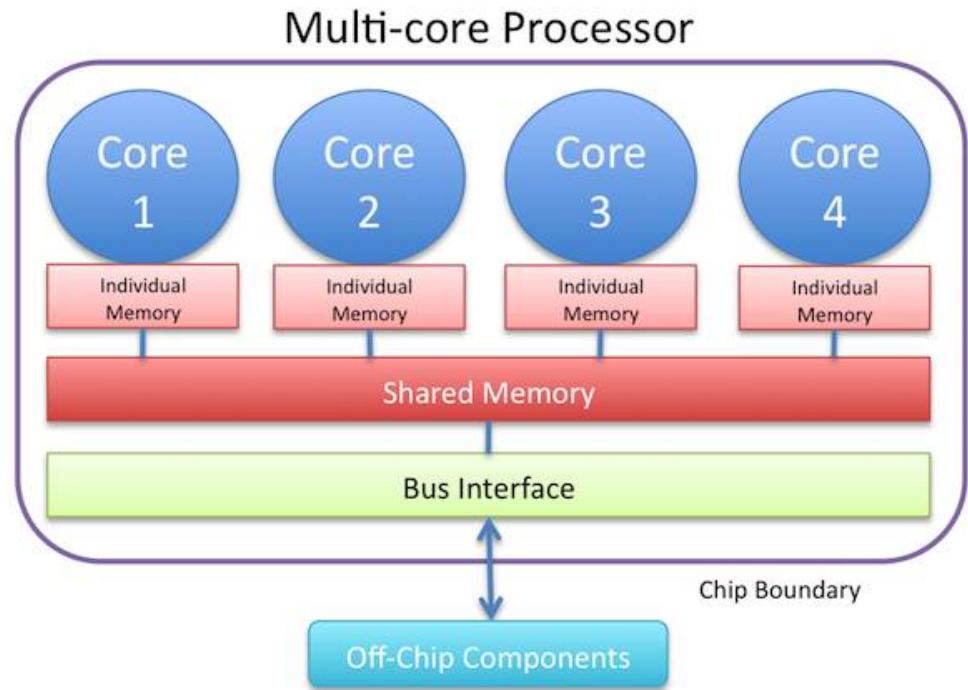


**4th Generation Intel® Core™ Processor Die Map**  
*22nm Tri-Gate 3-D Transistors*



## Modern CPU are multicore

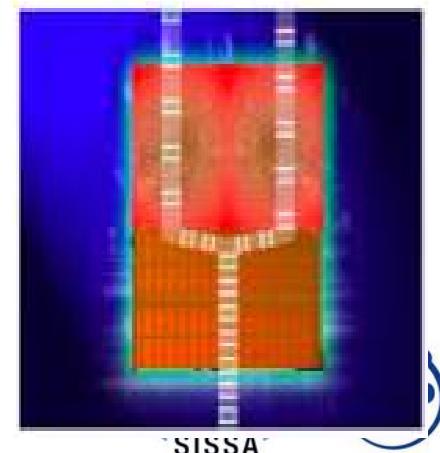
- Because of power, heat dissipation, etc increasing tendency to actually lower clock frequency but pack more computing cores onto a chip.
- These cores will share some resources, e.g. memory, network, disk, etc but are still capable of independent calculations



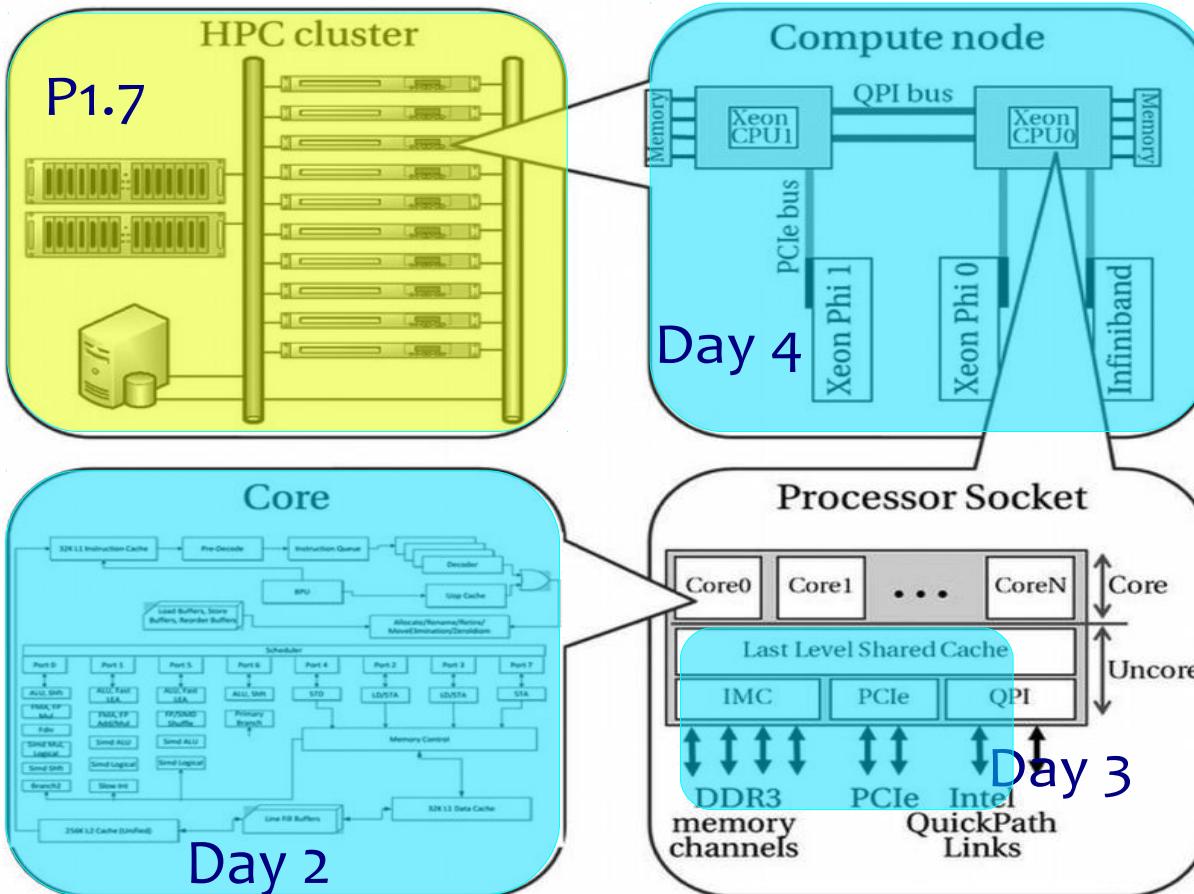
Picture from: <http://www.cse.wustl.edu/~jain/cse567-11/ftp/multcore/>

## What are cpu multi-core processors?

- Integrated circuit (IC) chips (CPU) containing more than one identical physical processor (core) . OS perceives each core as a discrete processor.
- Each core has its own complete set of resources, and may share the on-die cache layers (generally)
- Cores may have on-die communication path to front-side bus (FSB)
- What is a multi processor?
  - a collection of multicore cpus !



# First steps of the journey: core

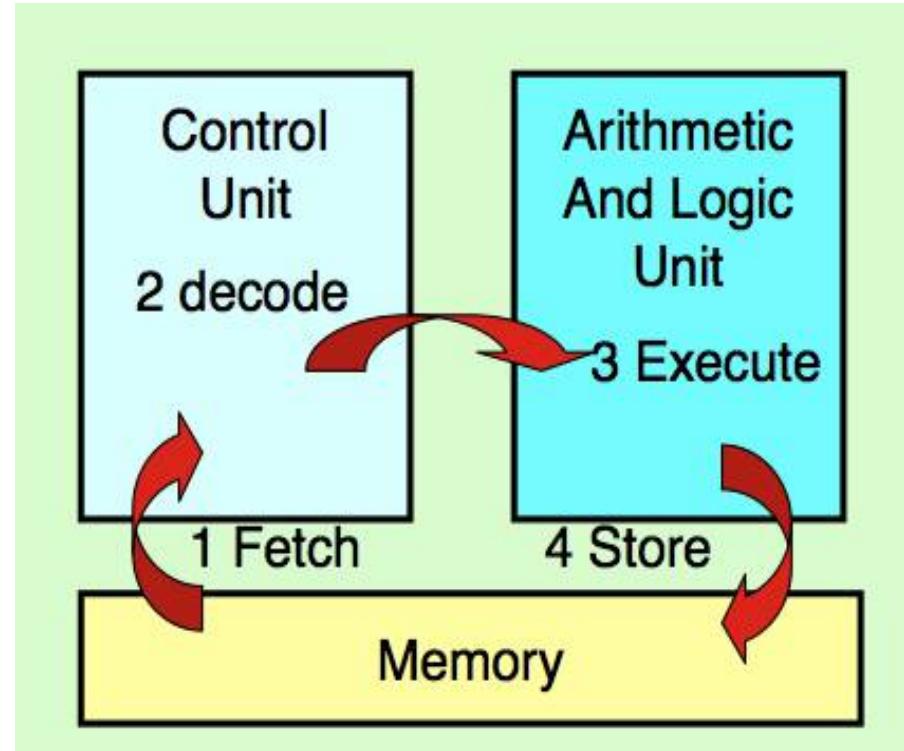


## What is a core?

Level 0 approximation:  
von Nuemann architecture

## Von Nuemann architecture:

- Control Unit: processes instructions ALU: math and logic operations
- At each cycle the CPU fetches both data and a description of what operations need to be performed and stores them in registers.



## Instruction Set Architecture (ISA)

- It is the boundary between SW and HW
- The deeper level accessible to the programmers
- The interface between the programmer and the microarchitecture
- Different microarchitectures can have the same ISA (binary Compatible)
- Different generation of microarchitectures can be backward compatible

## Analysis of a simple program on Von Neumann architecture:

```
void store(double *a, double *b,  
double *c){  
    *c = *a + *b;  
}
```

```
[exact@master ~]$ gcc -O2 -S -o - frammento.c  
    .file "frammento.c"  
    .text  
    .p2align 4,,15  
.globl store  
    .type    store, @function  
store:  
.LFB0:  
    .cfi_startproc  
    movsd  (%rdi), %xmm0  #load *a to mmx0  
    addsd  (%rsi), %xmm0  # load b and add to *a  
    movsd  %xmm0, (%rdx) # store to C  
    ret  
    .cfi_endproc  
.LFE0:  
    .size    store, .-store  
    .ident  "GCC: (GNU) 4.4.7 20120313 (Red Hat 4.4.7-4)"  
.section .note.GNU-stack,"",@progbits
```

## Workflow:

- Instruction fetch
- Instruction decode: determine operation and operands
- Memory fetch: Get operands from memory
- Perform operation
- Write results back
- Continue with next instruction

## How to go faster ?

- since many years CPU designers have been trying to improve and optimize the “Classical Model”:
  - increasing clock frequency (Moore’s law)
  - reducing the number of cycles needed to perform a single instruction ( $a = c + d$ )
    - flow/branch prediction
  - improve the number of outputs x cycle
    - Instructions Level Parallelism (ILP)
    - Execution => in-order Vs. out-of-order



**Contemporary cores are a little bit more complicated..**



## What is in a core really ?

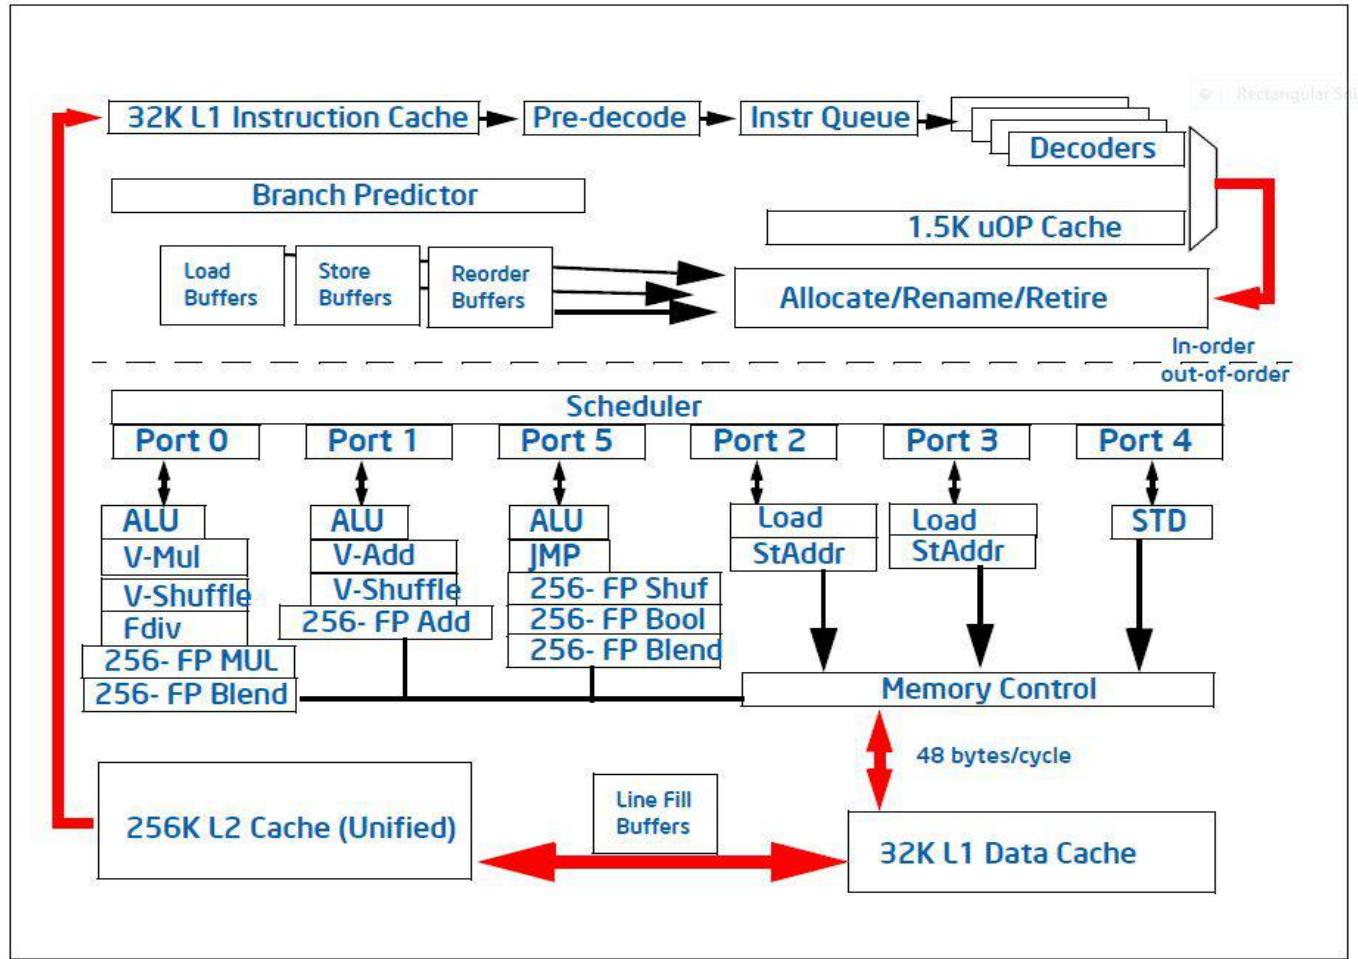


Figure 2-1. Intel microarchitecture code name Sandy Bridge Pipeline Functionality

## Instruction handling: out of order execution.

- It is unrealistic on modern CPU to assume all the instruction are executed sequentially
- “out of order” instruction handling has been present for almost 20 years
- Processor/core is allowed to change the order of the instructions in your code.
  - Be careful on this on Floating Point operations
- Several units on modern core are dedicated to this

## Instruction handling sections on Sandybridge core:

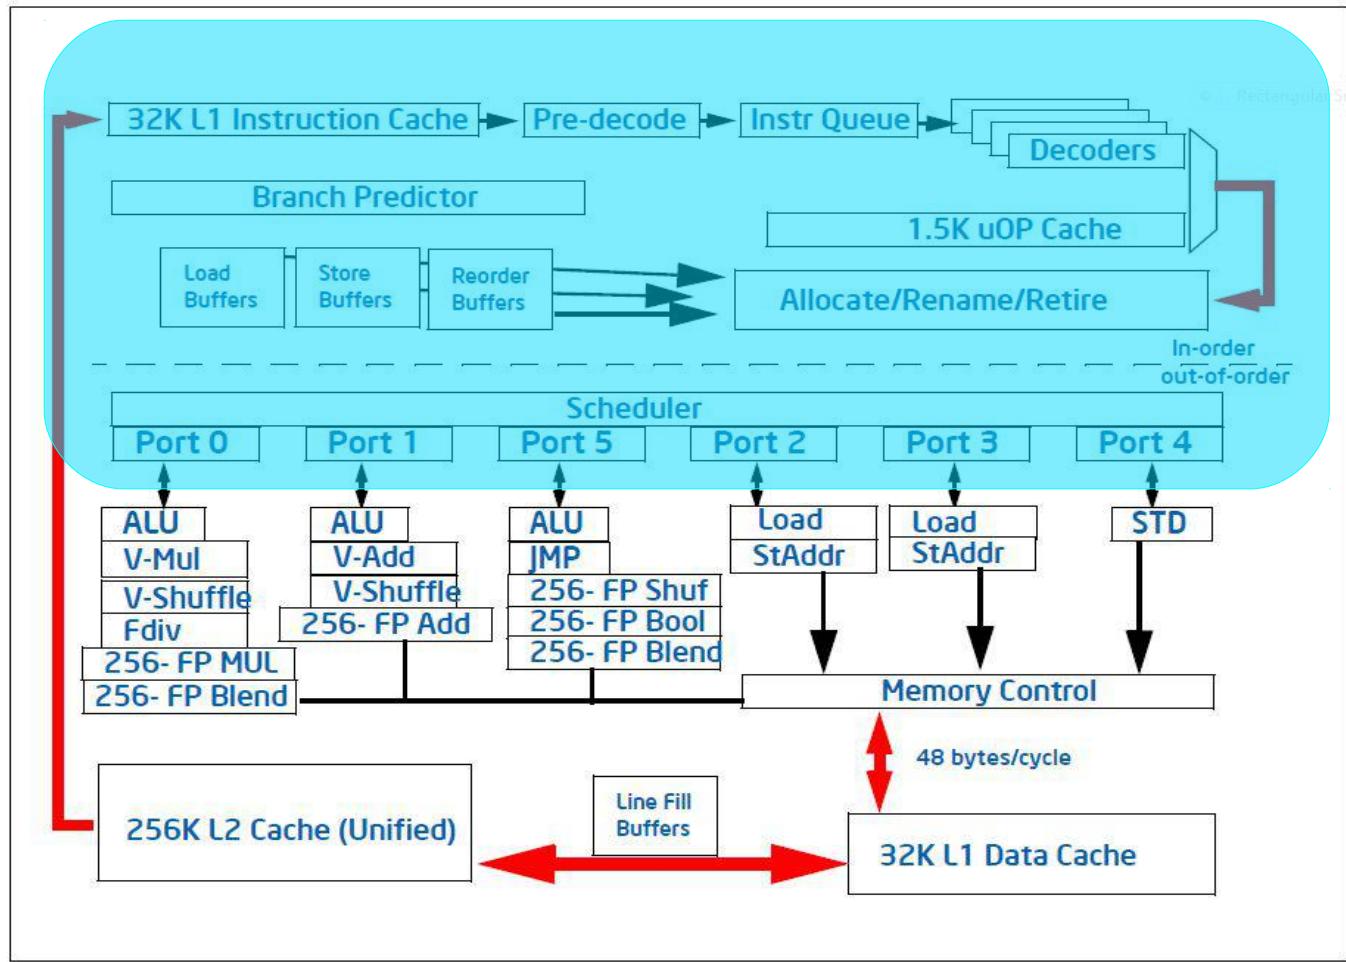


Figure 2-1. Intel microarchitecture code name Sandy Bridge Pipeline Functionality

## Many functional unit

- Circuitry on the chip which performs a given type of operation on operands in registers is known as a **functional unit**.
- For HPC we are mainly interested in Floating Point Units
- Some FPUs can perform more than one operation per cycle..
  - Multiple floating point units
  - 1 Mul + 1 Add, or 1 Fused Multiply-Add (FMA)

$$x \leftarrow c^*x+y$$

## What else in the cores ?

- On modern CPU/Cores there are many other stuff:
  - Superscalar execution
  - Pipelined functional units
  - Floating point instruction set extensions

## Superscalar CPU architecture

- The ability to issue/execute more than one instruction at the same time
- aka **Instruction level parallelism (ILP)**
- On sandybridge:
  - Six issue ports/ execution units that can execute Instruction simultaneously
  - e.g. Floating Point-Add (Port 1) and Floating-Point Mult (Port 0)

## Superscalarity on Sandybridge

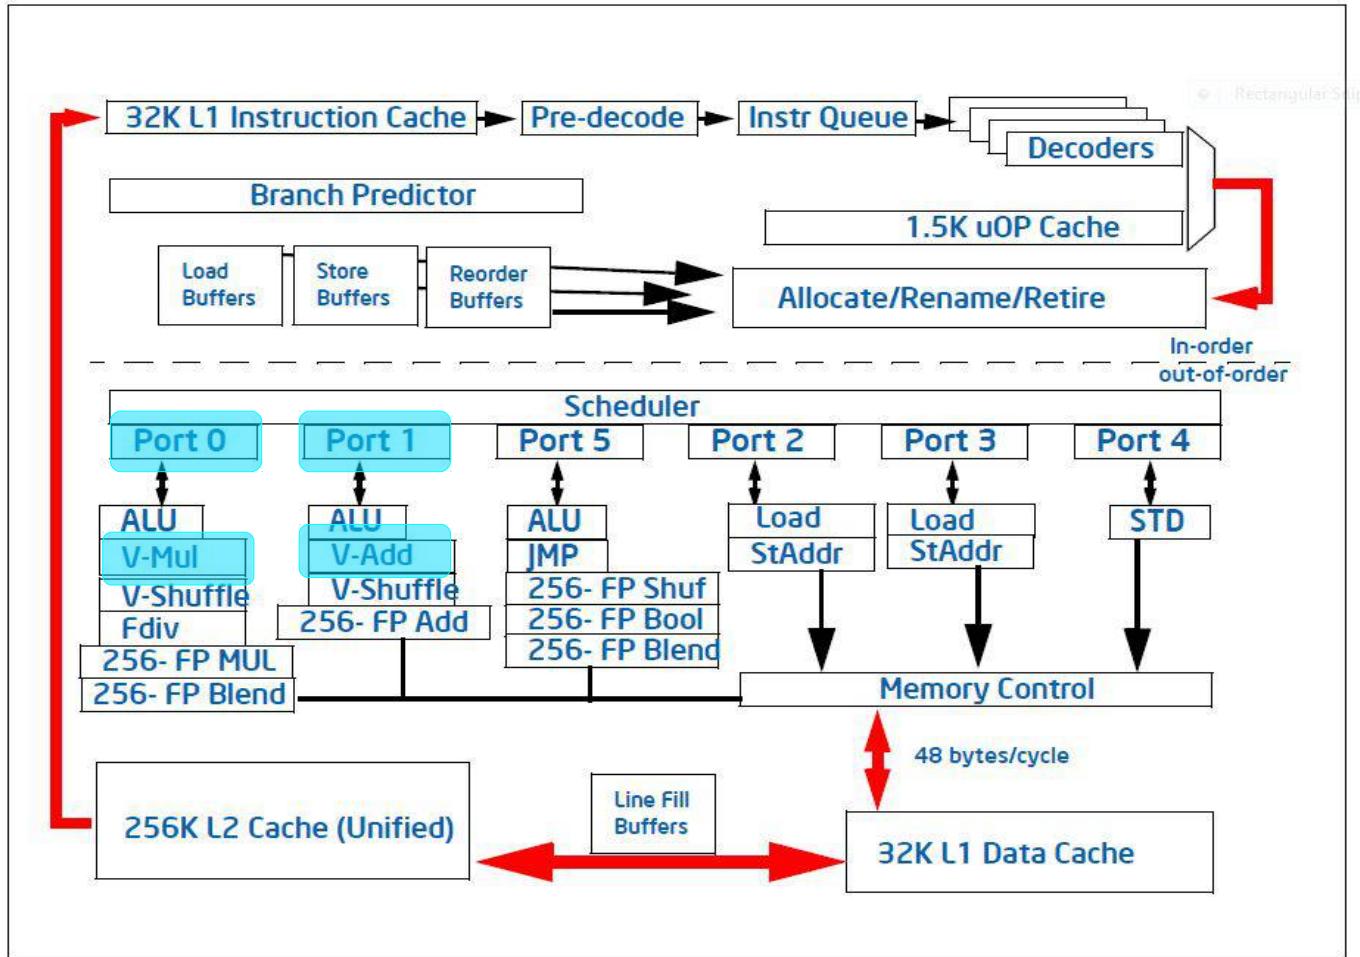


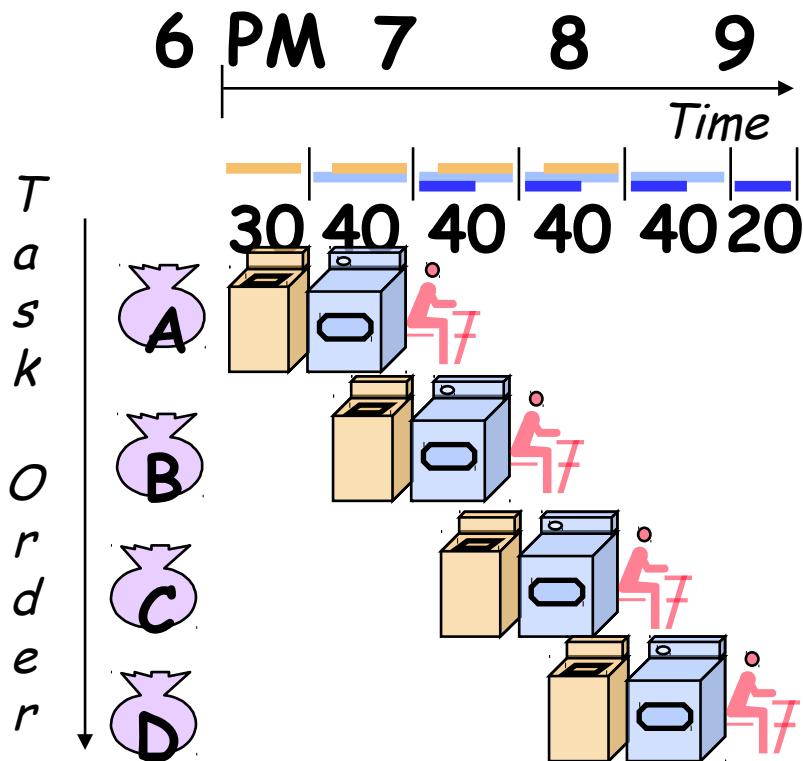
Figure 2-1. Intel microarchitecture code name Sandy Bridge Pipeline Functionality

## Pipelined Functional Units

- Most integer and floating point functional units are pipelined, meaning that they can have multiple independent executions of the same instruction placed in a queue.
- The idea is that after an initial startup latency, the functional unit should be able to generate one result every clock period (CP).
- Each stage of a pipelined operation can be working simultaneously on different sets of operands.
- For each stage of the pipeline there is a dedicated HW

## What is Pipelining?

Dave Patterson's Laundry example: 4 people doing laundry  
wash (30 min) + dry (40 min) + fold (20 min)

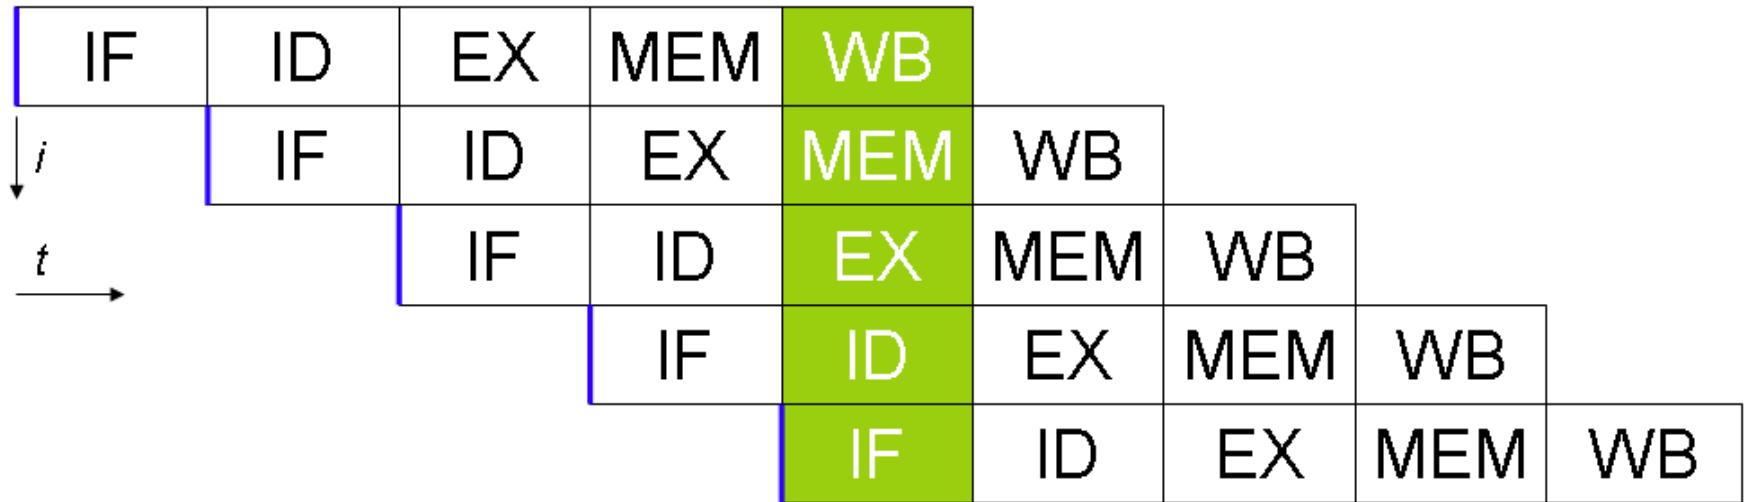


In this example:

Sequential execution  
takes  $4 * 90\text{min} = 6$   
hours

Pipelined execution  
takes  $30 + 4 * 40 + 20 = 3.3$   
hours

## The standard RISC 5 stage pipeline



IF = Instruction Fetch

ID = Instruction Decode

EX = Execute,

MEM = Memory access

WB = Register write back

## Pipeline concepts/jargon

- Wind-up phase: time it takes to load the pipeline (a.k.a. latency)
- Wind-down phase: time it takes to drain the pipeline
- Segments=stages=steps of the pipeline
- Speedup:  $T_{seq}/T_{pipe}$
- Throughput =  $N/T_{pipe}$
- Pipelining helps throughput, but not latency
- Pipeline rate limited by slowest pipeline stage

## Pipeline analysis: $N_{1/2}$

Given a pipeline with  $s$  segments and  $N$  operations:

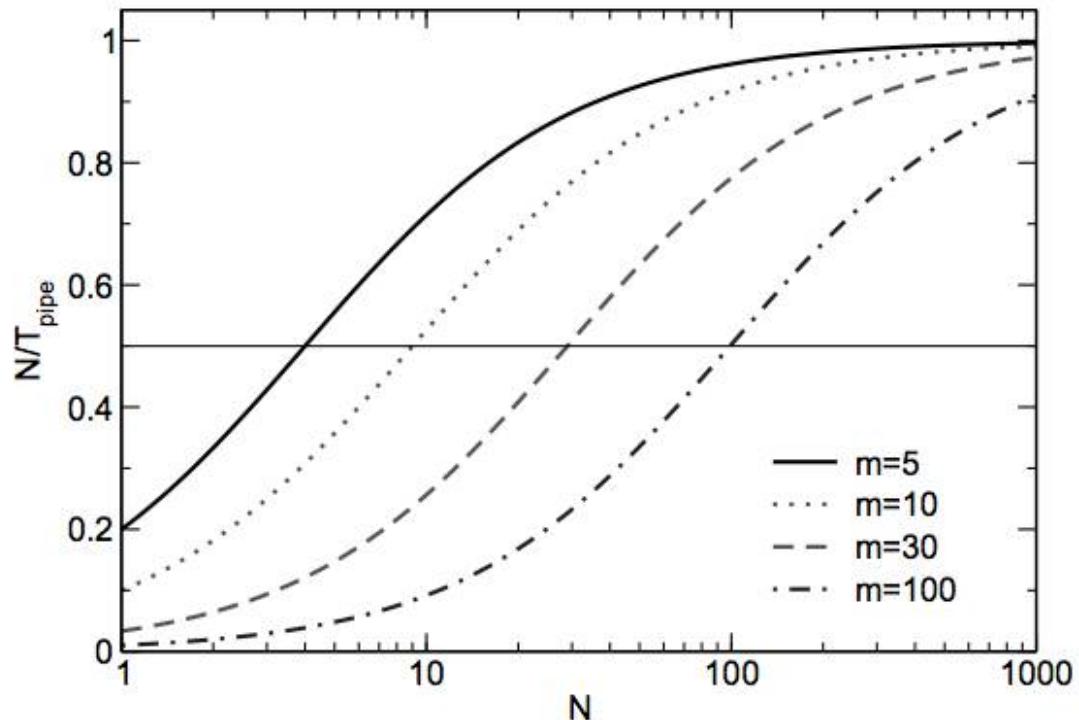
- $T_{\text{seq}}$  (the time without pipelining) =  $sN$
- $T_{\text{pipe}} = s+N-1$

$$\rightarrow \text{Speed-up} = s*N/(s+N-1)$$

$$\text{Throughput} = N/T_{\text{pipe}} = 1 / (1 + (s-1)/N)$$

- for  $N$  large enough
  - Speedup  $\sim s$
  - Throughput  $\sim 1$ 
    - i.e. 1 result per clock cycle
- With  $N$  operations, actual rate is  $N/(s+N)$
- This is half of the asymptotic rate if  $s=N$

## Asymptotic behavior of pipeline throughput

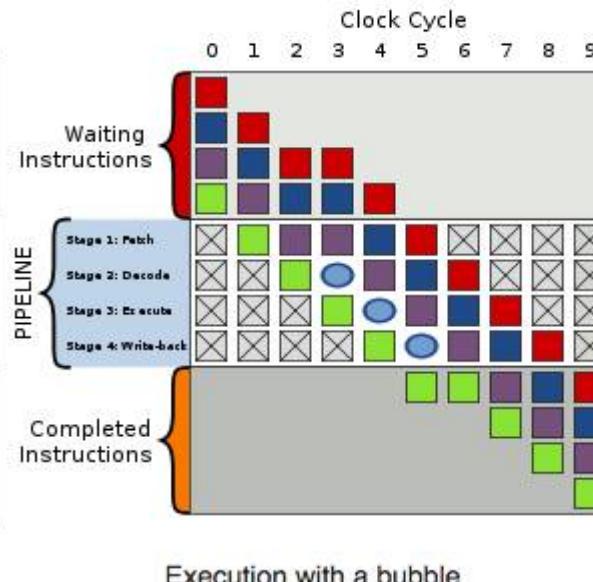
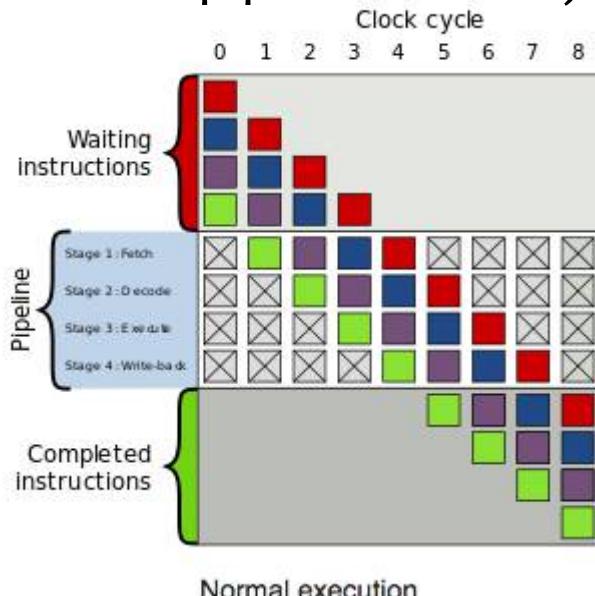


**Figure 1.6:** Pipeline throughput as a function of the number of independent operations.  $m$  is the pipeline depth.

## Pipeline drawbacks

Standard lengths of modern processor: from 10 to 35

- Not efficient at all if loops are short and tight
- Not efficient if there are jumps (branches) or stalling in instructions (the so called pipeline bubble)



## Branch Prediction

- The “instruction pipeline” is all of the processing steps (also called segments) that an instruction must pass through to be “executed”.
- Higher frequency machines have a larger number of segments.
- Branches are points in the instruction stream where the execution may jump to another location, instead of executing the next instruction.
- For repeated branch points (within loops), instead of waiting for the loop to branch route outcome, it is predicted.

## Floating Point Instruction Set Extensions

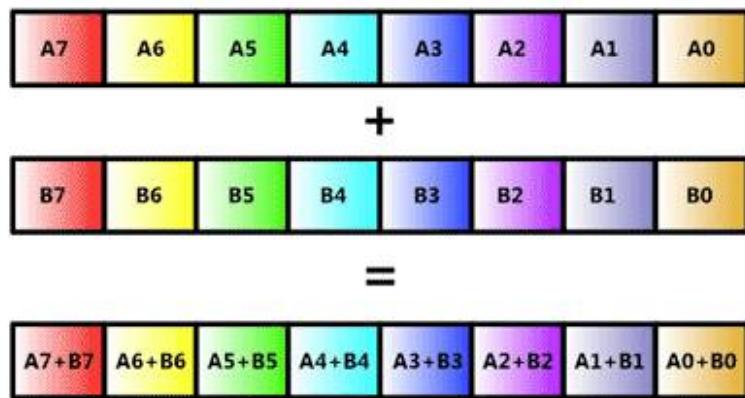
- additional floating point instructions beyond the usual floating point add and multiply instructions:
  - Square root instruction --usually not pipelined!
    - AMD Opteron / Intel Xeon
  - SIMD (a.k.a. vector) floating point instructions
    - AMD Opteron/ Intel Xeon
- Combined floating point multiply/add (MADD) instruction
  - AMD Opteron ("Barcelona" and after, using SIMD)
  - Intel Xeon ("Woodcrest" and after, using SIMD)

## Instruction Set Extensions

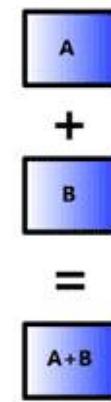
- Intel
  - MMX (Matrix Math eXtensions)
  - SSE (Streaming SIMD Extensions)
  - SSE2 / SSE4.2
  - AVX AVX-512 (Advanced vector eXtension)
    - From sandybridge processor on
- AMD
  - 3DNow!
  - AMD 3DNow!+ (or 3DNow! Professional, or 3DNow! Athlon) ...
- To check what you have on your machine:
  - cat /proc/cpuinfo
  - to enable them: use appropriate compiler flag..

## SIMD approach:

### SIMD Mode



### Scalar Mode



## The AVX Instruction Set

Latest version of Intel and GNU compiler are in most cases able to perform auto-vectorization on simple loop constructs and with simplified data-structures

Machine-dependent vector intrinsic that use a function call syntax to emit machine code (to be used at large when compilers not work as expected)

Memory alignment is required

It is available as compiler option on both Intel and GNU compiler:

-mavx (GNU) | -xAVX (INTEL)

## Intel SIMD evolution

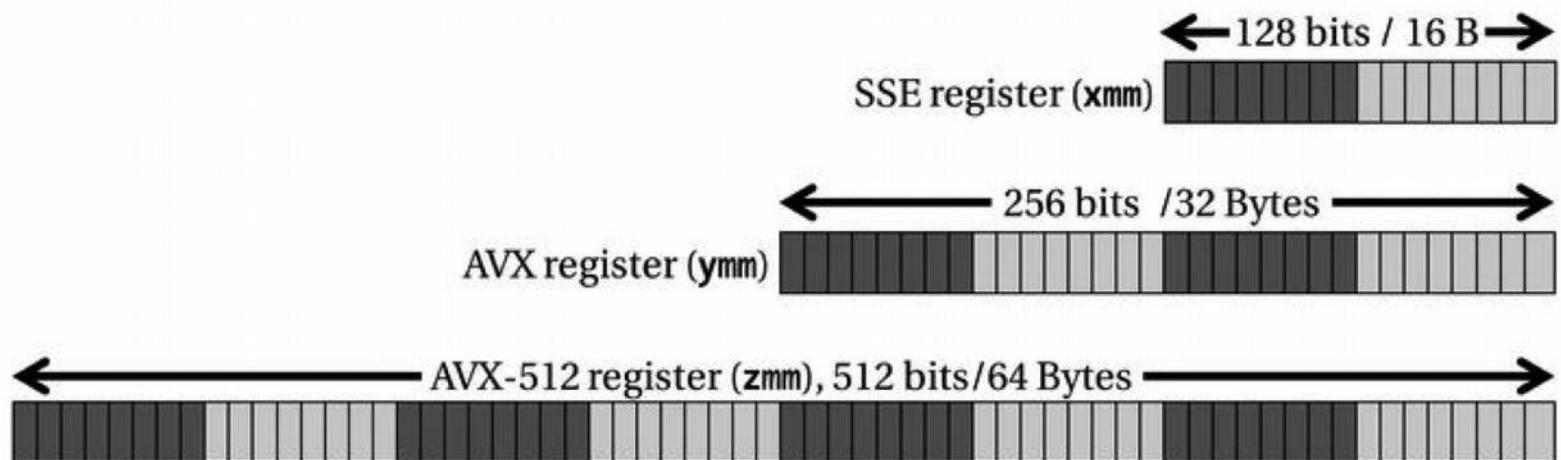


Figure 2-9. SSE, AVX, and AVX-512 vector registers with packed 64-bit numbers

## Number crunching on CPU: what do we count ?

- Rate of [million/billions of] floating point operations per second ([M|G]flops) FLOPs/S
- **Theoretical peak performance:**

determined by counting the number of floating-point additions and multiplications that can be completed during a period of time, usually the cycle time of the machine

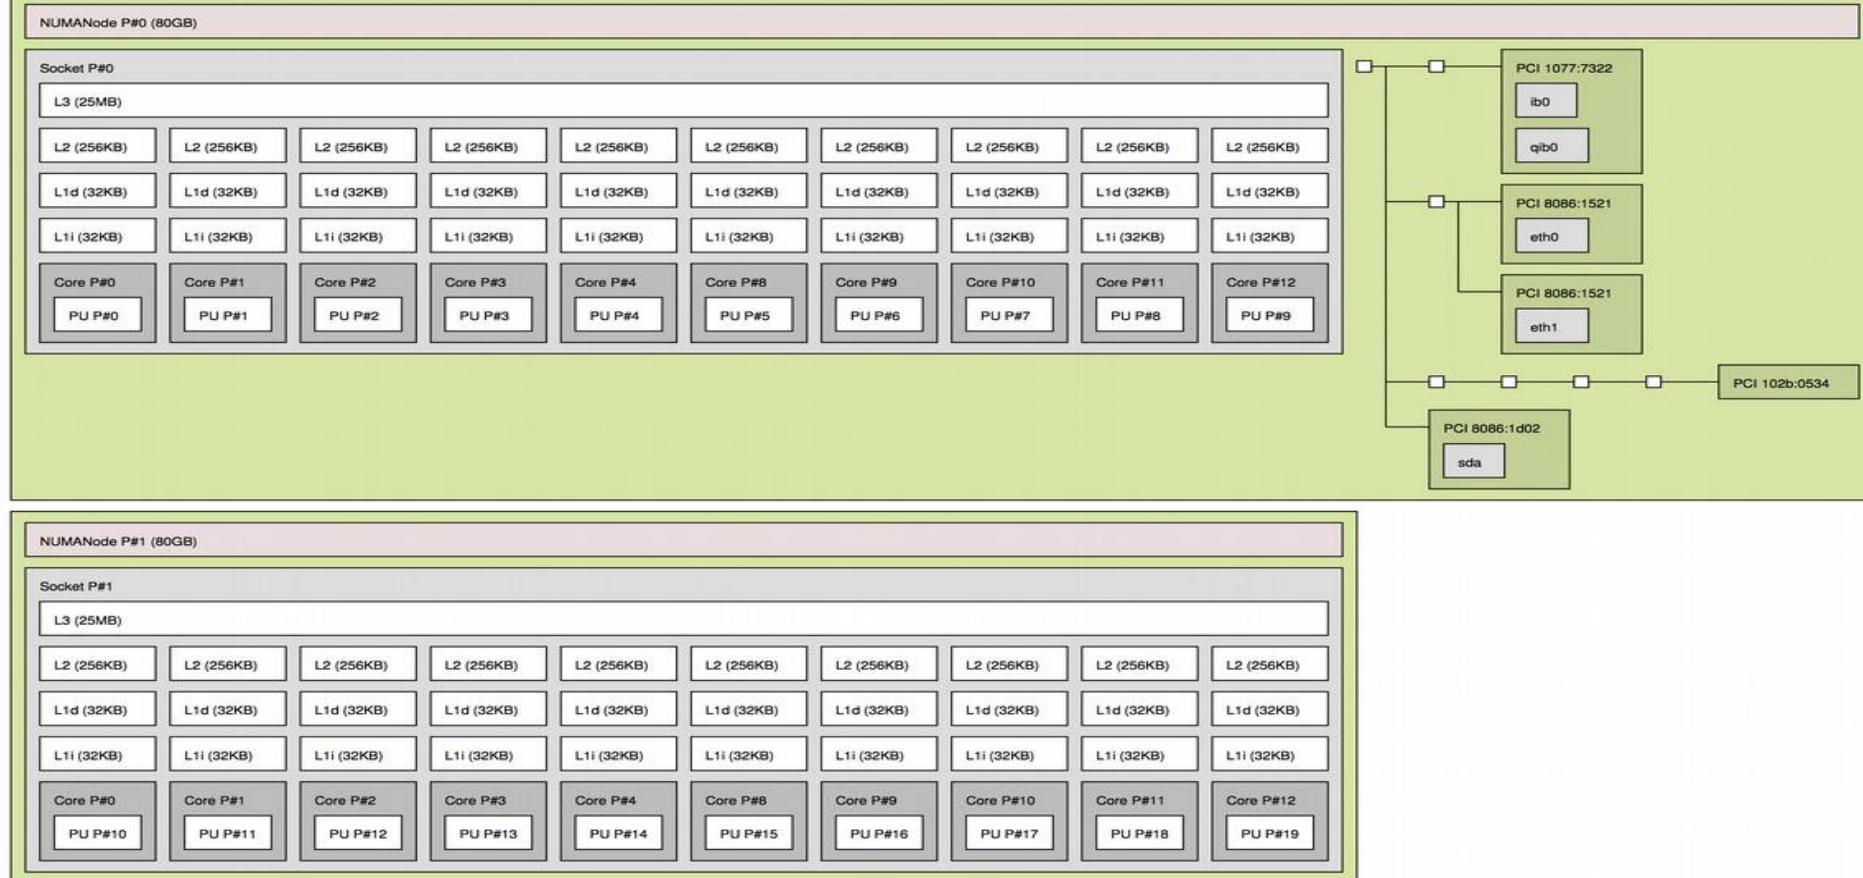
$$\text{FLOPS} = \text{Clock-rate} * \text{Number\_of\_FP\_operation} * \text{Number\_of\_cores}$$

## Exercise: compute Theoretical Peak performance for your laptop/desktop:

- Identify the CPU
- Identify the frequency
- Identify the number of floating point for cycle
- Identify how many cores
- Put all together in one single number

# A modern node picture (Ulysses-node)

Machine (160GB)



Host: cn08-19

Indexes: physical

Date: Wed Sep 16 11:06:52 2015

## Peak performance on ulysses node:

- Two Sockets
- 10 Cores per Socket @ 2.6 GHz
- Two-way superscalar with respect to floating-point
  - AVX Mult on Port 0, AVX Add on Port 1
- 256b AVX
  - Four double-precision operations with a single instruction

$2 \times 10 \times 2.6 \text{ GHz} \times 2 \text{ instr/cycle} \times 4 \text{ ops/instr} = \text{????? Gflop/s}$

# What CPU for HPC?

Processor Generation	Count	System Share (%)	Rmax (GFlops)	Rpeak (GFlops)	Cores
Intel Xeon E5 (Haswell)	221	44.2	172,605,096	271,881,324	8,636,254
Intel Xeon E5 (IvyBridge)	118	23.6	109,146,786	187,108,623	8,639,336
Intel Xeon E5 (SandyBridge)	66	13.2	48,003,244	67,528,012	3,222,875
Intel Xeon E5 (Broadwell)	21	4.2	28,625,858	33,317,165	940,140
Power BQC	19	3.8	47,047,481	55,155,100	4,308,992
Xeon 5600-series (Westmere-EP)	9	1.8	9,662,361	19,393,022	641,068
Intel Xeon E7 (IvyBridge)	8	1.6	3,291,391	6,954,496	398,880
Intel Xeon E7 (Haswell-Ex)	7	1.4	3,429,884	5,678,592	159,456
Opteron 6300 Series "Abu Dhabi"	5	1	1,894,061	6,710,298	656,128
SPARC64 XIIfx	5	1	10,422,200	11,530,310	354,384
Opteron 6200 Series "Interlagos"	5	1	20,216,100	30,720,688	862,264
POWER7	4	0.8	2,788,135	3,337,539	112,448
Opteron 6100-series "Magny-Cours"	3	0.6	1,704,800	2,354,995	232,704
Xeon 5500-series (Nehalem-EP)	2	0.4	930,000	1,510,046	71,880
Xeon E7 (Westmere-EX)	1	0.2	354,300	655,258	38,400
ShenWei	1	0.2	795,900	1,070,160	137,200
Intel Xeon Phi	1	0.2	817,847	1,474,458	32,912
SPARC64 IXfx	1	0.2	1,043,000	1,135,411	76,800
Xeon 5500-series (Nehalem-EX)	1	0.2	1,050,000	1,254,550	138,368
SPARC64 VIIIfx	1	0.2	10,510,000	11,280,384	705,024
Sunway	1	0.2	93,014,594	125,435,904	10,649,600

## Some considerations

- Intel is dominating
- AMD is disappearing (or already disappeared)
- IBM is coming back with Power8 architecture
- ARM processor still to come

## Intel Xeon history...

The Process Gap Between PCs And Servers At Intel

<b>Chip Family</b>	<b>Process</b>	<b>PC Chip</b>	<b>Launch</b>	<b>Server Chip</b>	<b>Launch</b>	<b>Months</b>
			<b>Date</b>		<b>Date</b>	
Nehalem	45 nm	Core i7	Nov-08	Xeon 5500	Mar-09	4
Westmere	32 nm	Core i7	Jul-10	Xeon 5600	Feb-11	7
Sandy Bridge	32 nm	Core i7	Jan-11	Xeon E5	Mar-12	14
Ivy Bridge	22 nm	Core i7	Apr-12	Xeon E5 v2	Sep-13	17
Haswell	22 nm	Core i7	Jun-13	Xeon E5 v3	Sep-14	15
Broadwell	14 nm	Core i7	Jan-15	Xeon E5 v4	<b>Mar-16</b>	<b>14</b>
Skylake	14 nm	Core i7	<b>Aug-15</b>	Xeon E5 v5	<b>Mar-17</b>	<b>19</b>
Kaby Lake	14 nm	Core i7	<b>Aug-16</b>	Xeon E5 v6	<b>Mar-18</b>	<b>19</b>
Cannonlake	10 nm	Core i7	<b>Aug-17</b>	Xeon E5 v7	<b>Mar-19</b>	<b>19</b>
???	10 nm	Core i7	<b>Aug-18</b>	Xeon E5 v8	<b>Mar-20</b>	<b>19</b>
???	10 nm	Core i7	<b>Aug-19</b>	Xeon E5 v9	<b>Mar-21</b>	<b>19</b>
???	7 nm	Core i7	<b>Aug-20</b>	Xeon E5 v10	<b>Mar-22</b>	<b>19</b>
???	7 nm	Core i7	<b>Aug-21</b>	Xeon E5 v11	<b>Mar-23</b>	<b>19</b>
???	7 nm	Core i7	<b>Aug-22</b>	Xeon E5 v12	<b>Mar-24</b>	<b>19</b>

Source: *The Platform*

## Intel processor families:

### High Core Count Top Bin

<b>Family</b>	<b>Model</b>	<b>Clock Speed</b>	<b>Cores / Threads</b>	<b>1K Tray Unit Price</b>	<b>Rel Perf</b>	<b>\$ / Rel Perf</b>
Nehalem	X5570	2.93 GHz	4 / 8	\$1,386	116	\$1,195
Westmere	X5680	3.33 GHz	6 / 12	\$1,663	198	\$840
Sandy Bridge	E5-2690	2.9 GHz	8 / 16	\$2,057	2.55	\$806
Ivy Bridge	E5-2697 v2	2.7 GHz	12 / 24	\$2,614	3.73	\$702
Haswell	E5-2699 v3	2.3 GHz	18 / 36	\$4,115	5.20	\$792
Broadwell	E5-2699 v4	2.2 GHz	22 / 44	\$4,115	6.34	\$649

### Standard Top Bin

<b>Family</b>	<b>Model</b>	<b>Clock Speed</b>	<b>Cores / Threads</b>	<b>1K Tray Unit Price</b>	<b>Rel Perf</b>	<b>\$ / Rel Perf</b>
Nehalem	E5540	2.53 GHz	4 / 8	\$744	100	\$744
Westmere	E5640	2.66 GHz	4 / 8	\$744	104	\$744
Sandy Bridge	E5-2640	2.5 GHz	6 / 12	\$885	165	\$536
Ivy Bridge	E5-2640 v2	2.0 GHz	8 / 16	\$885	184	\$481
Haswell	E5-2640 v3	2.6 GHz	8 / 16	\$940	2.61	\$360
Broadwell	E5-2640 v4	2.4 GHz	10 / 20	\$939	3.14	\$299

### Highest Clock Speed

<b>Family</b>	<b>Model</b>	<b>Clock Speed</b>	<b>Cores / Threads</b>	<b>1K Tray Unit Price</b>	<b>Rel Perf</b>	<b>\$ / Rel Perf</b>
Nehalem	X5570	2.93 GHz	4 / 8	\$1,386	116	\$1,195
Westmere	X5677	3.46 GHz	4 / 8	\$1,663	136	\$1,223
Sandy Bridge	E5-2643	3.3 GHz	4 / 8	\$885	145	\$610
Ivy Bridge	E5-2637 v2	3.5 GHz	4 / 8	\$996	1.61	\$619
Haswell	E5-2637 v3	3.5 GHz	4 / 8	\$995	176	\$566
Broadwell	E5-2637 v4	3.5 GHz	4 / 8	\$996	1.83	\$543

## Intel® Xeon® Processor E5 v4 Product Family Overview

### New features:

- Broadwell microarchitecture
- Built on 14nm process technology
- Socket compatible<sup>◊</sup> replacement for Intel<sup>®</sup> Xeon<sup>®</sup> processor E5-2600 v3 on Grantley

### New processor technologies:

- Posted Interrupts
- Page Modification Logging
- Cache Allocation Technology
- Memory BW Monitoring

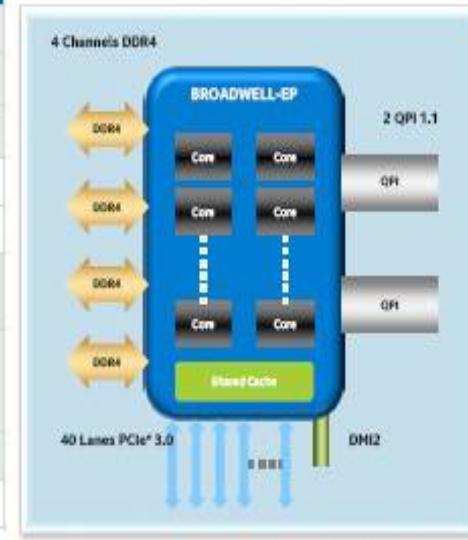
- Crypto Speedup
- Supervisor Mode Access Prevention
- New RDSEED instruction
- Intel<sup>®</sup> Processor Trace
- Hardware Controlled Power Management

Features	Xeon E5-2600 v3 (Haswell-EP)	Xeon E5-2600 v4 (Broadwell-EP)
Cores Per Socket	Up to 18	Up to 22
Threads Per Socket	Up to 36 threads	Up to 44 threads
Last-level Cache (LLC)	Up to 45 MB	Up to 55 MB
QPI Speed (GT/s)	2x QPI 1.1 channels 6.4, 8.0, 9.6 GT/s	
PCIe <sup>*</sup> Lanes / Speed(GT/s)	40 / 10 / PCIe <sup>*</sup> 3.0 (2.5, 5, 8 GT/s)	
Memory Population	4 channels of up to 3 RDIMMs or 3 LRDIMMs	+ 3DS LRDIMM <sup>†</sup>
Memory RAS	ECC, Patrol Scrubbing, Demand Scrubbing, Sparing, Mirroring, Lockstep Mode, x4/x8 SDDC	+ DDR4 Write CRC
Max Memory Speed	Up to 2133	Up to 2400
TDP (W)	160 (Workstation only), 145, 135, 120, 105, 90, 85, 65, 55	

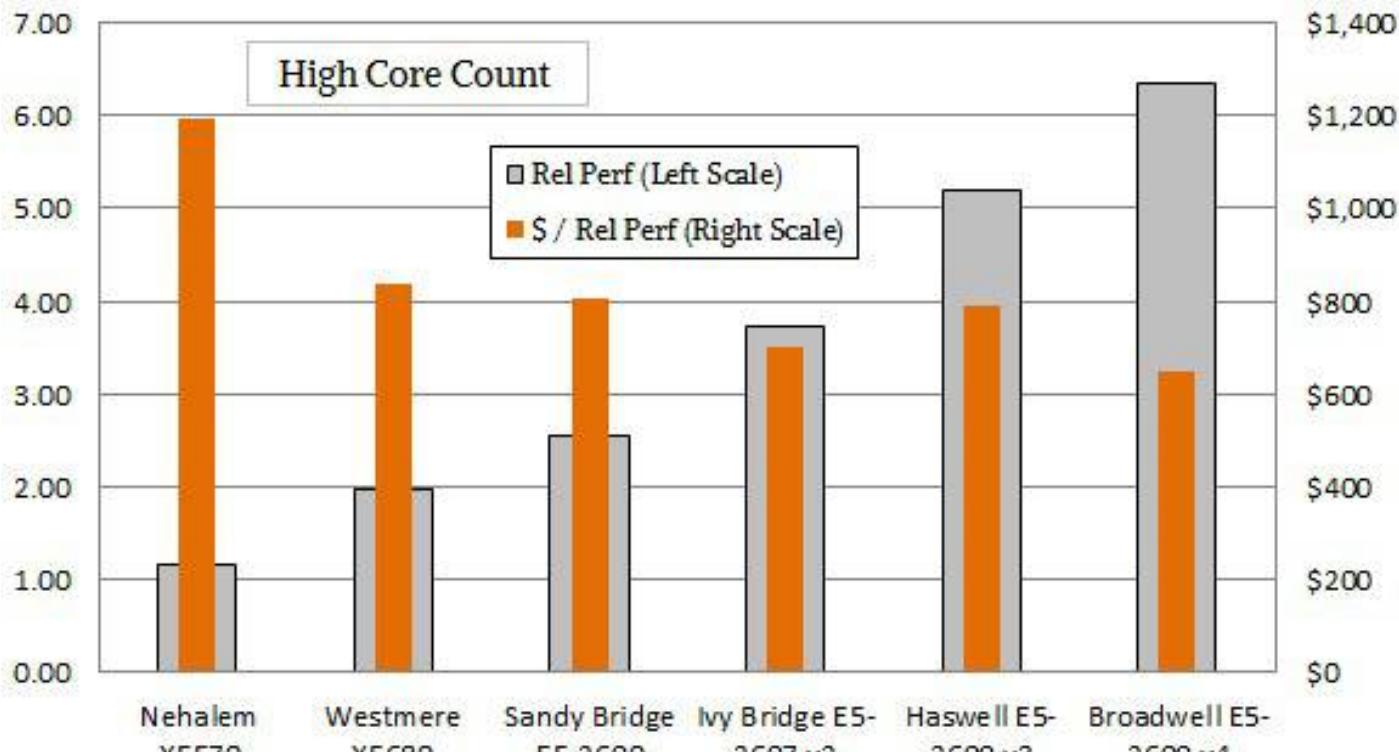
<sup>◊</sup> Requires BIOS and firmware update

<sup>†</sup> Depends on market availability

All products, computer systems, dates and figures specified are preliminary based on current expectations, and are subject to change without notice. Intel may make changes to specifications and product descriptions at any time, without notice.

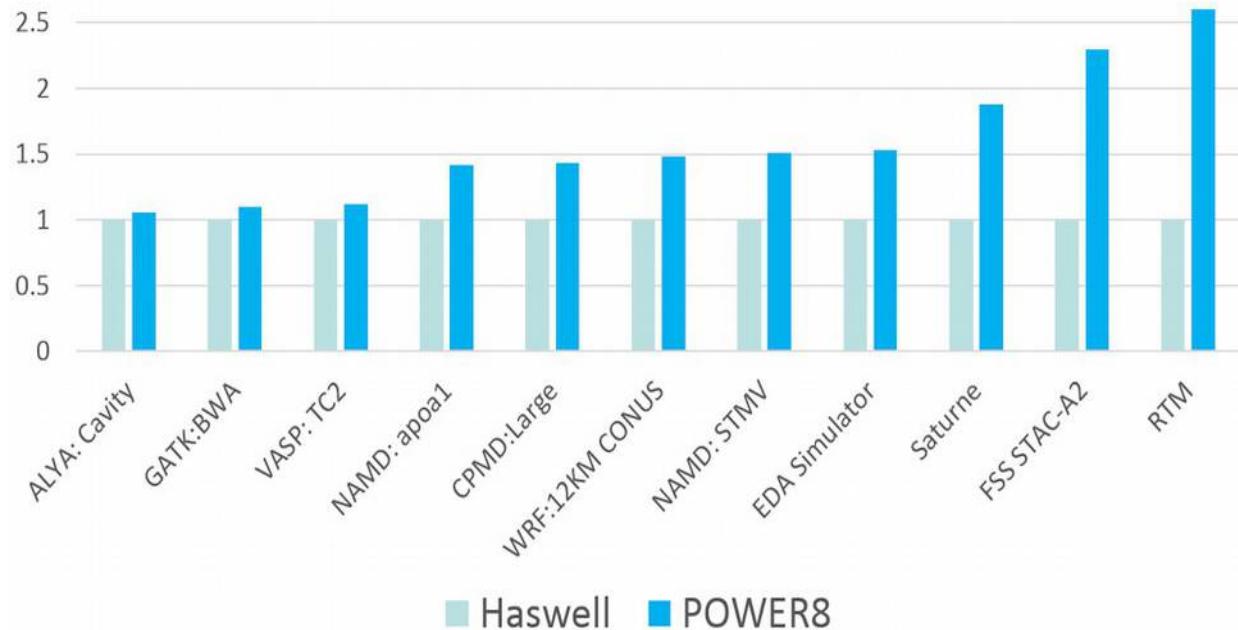


## Performance vs \$/performance



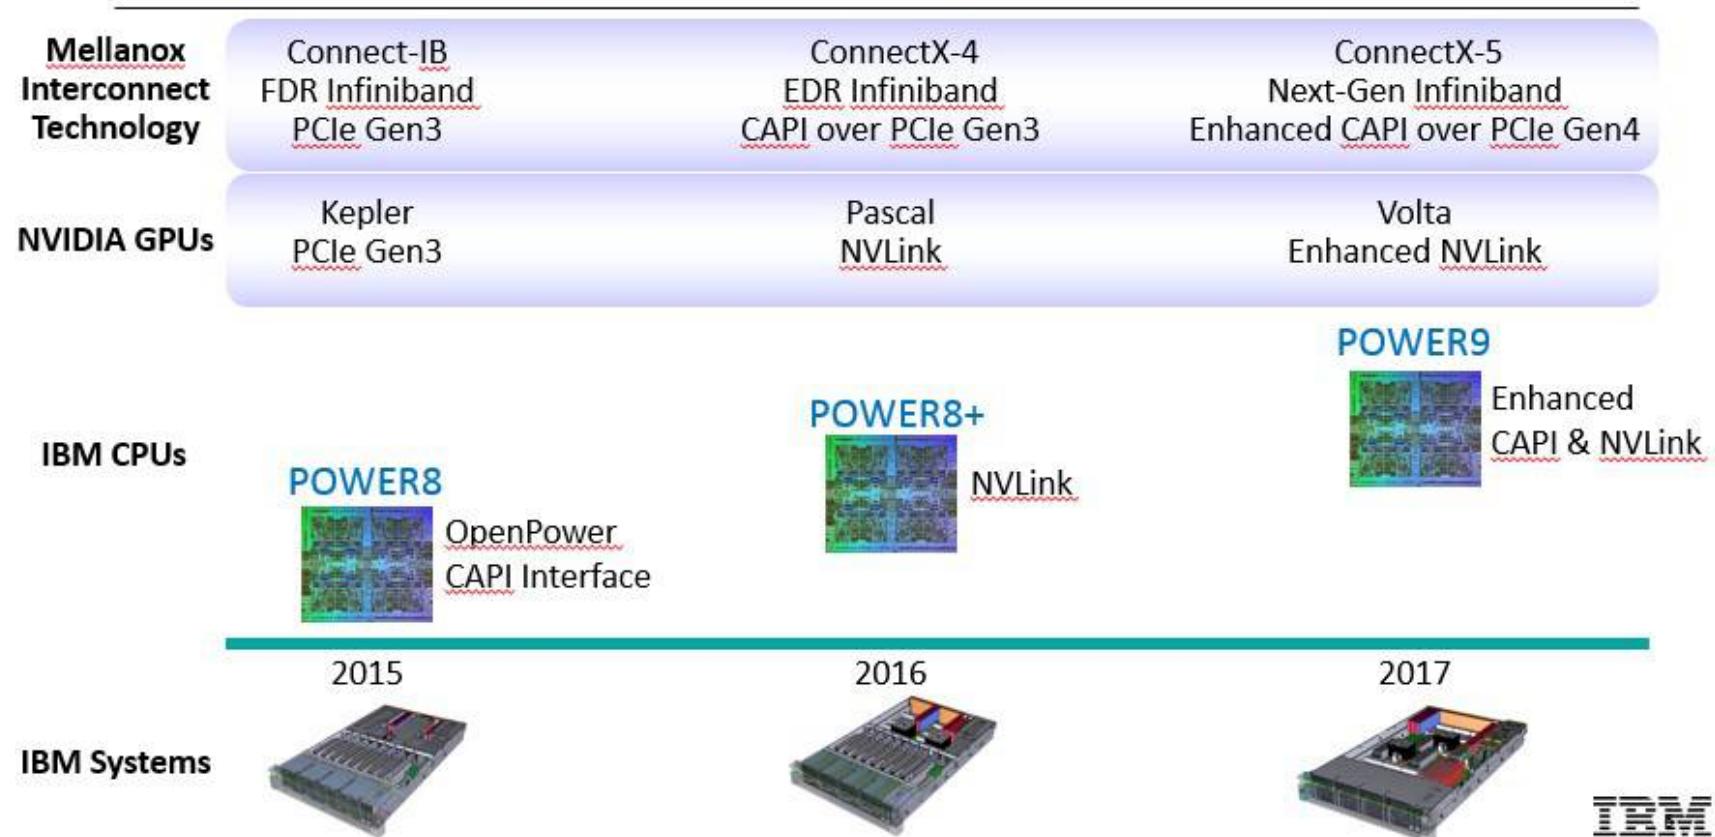
Source: *The Next Platform*

## Power8 vs Intel Xeon



From: <http://www.nextplatform.com/2015/07/15/ibm-readies-power8-for-openpower-push/>

## Power8 roadmap



## Conclusions

- Modern CPU have **a high degree of parallelism** some time hidden to the user
- In order to optimize on them you should be aware of this.
- Next presentation: some trick of the trade to understand bottleneck and increase performance